

SUPPLEMENTARY MATERIAL

Throughout the following sections, we use the notation and references from the paper.

1 RAP MESSAGES

Application of the max-sum algorithm [32] gives two types of messages:

- (1) Message from variable to factor/constraint node is simply the sum of all the incoming messages to the variable node except the message from the target factor node. This can be observed in β_{ij} , ϕ_{ij} and ρ_{ij} messages in (18), (20) and (22) respectively.
- (2) Message from factor node to variable node depends upon the corresponding factor/constraint function. Since the $E(\cdot)$ constraint is the same in RAP as in AP, the messages α_{ij} in (23) are the same as in (9). Therefore we are left with the messages $\bar{\eta}_{ij}$ and ψ_{ij} that are derived below.

1.1 Derivation of $\bar{\eta}_{ij}$

According to max-sum algorithm:

$$\bar{\eta}_{ij}(h_{ij}) = \max_{k \neq j} \left(\bar{I}(H[i, :]) + \sum_{l \neq j} \beta_{il}(h_{il}) \right) \quad (27)$$

When $h_{ij} = 1$: We have met the least requirement of $\bar{I}(\cdot)$ constraint to have at least one entry in i th row equal to 1. For the rest of points q governs whether or not they will be chosen by x_i .

$$\begin{aligned} \bar{\eta}_{ij}(1) &= \max_{k \neq j} \left(\sum_m h_{im} q + \sum_{l \neq j} \beta_{il}(h_{il}) \right) \\ &= \max_{k \neq j} \left(\sum_{l \neq j} (q h_{il} + \beta_{il}(h_{il})) \right) + q \\ &= \sum_{k \neq j} \left(\max_{h_{ik} \in \{0, 1\}} (q h_{ik} + \beta_{ik}(h_{ik})) \right) + q \\ &= \sum_{k \neq j} \max (q + \beta_{ik}(1), \beta_{ik}(0)) + q \\ &= \sum_{k \neq j} \max (q + \beta_{ik}, 0) + \sum_{k \neq j} \beta_{ik}(0) + q \end{aligned} \quad (28)$$

When $h_{ij} = 0$: We need to ensure at least one $h_{ik} = 1; k \neq j$. The rest can decide for themselves based upon the q .

$$\begin{aligned} \bar{\eta}_{ij}(0) &= \max_{m \neq j} \beta_{im}(1) \\ &+ \max_{k \neq \{m, j\}} \left(\sum_{l \neq \{m, j\}} (q h_{il} + \beta_{il}(h_{il})) \right) + q \\ &= \max_{m \neq j} \beta_{im}(1) + \sum_{k \neq \{m, j\}} \left(\max_{h_{ik} \in \{0, 1\}} (q h_{ik} + \beta_{ik}(h_{ik})) \right) + q \\ &= \max_{m \neq j} \beta_{im}(1) + \sum_{k \neq \{m, j\}} \max (q + \beta_{ik}(1), \beta_{ik}(0)) + q \\ &= \max_{m \neq j} \beta_{im}(1) + \sum_{k \neq \{m, j\}} \max (q + \beta_{ik}, 0) \end{aligned}$$

$$+ \sum_{k \neq \{m, j\}} \beta_{ik}(0) + q \quad (29)$$

For binary variable h_{ij} , instead of propagating separate messages for $h_{ij} = 0$ and $h_{ij} = 1$, it is enough to propagate the difference [13]. Therefore the message $\bar{\eta}_{ij}$ is given by:

$$\begin{aligned} \bar{\eta}_{ij} &= \bar{\eta}_{ij}(1) - \bar{\eta}_{ij}(0) \\ &= \sum_{k \neq j} \max (q + \beta_{ik}, 0) + \sum_{k \neq j} \beta_{ik}(0) - \max_{m \neq j} \beta_{im}(1) \\ &- \sum_{k \neq \{m, j\}} \max (q + \beta_{ik}, 0) - \sum_{k \neq \{m, j\}} \beta_{ik}(0) \\ &= - \max_{m \neq j} (\beta_{im}(1) - \max (q + \beta_{im}, 0) - \beta_{im}(0)) \\ &= - \max_{m \neq j} (\beta_{im} + \min (-q - \beta_{im}, 0)) \\ &= - \max_{m \neq j} (\min (\beta_{im}, -q)) \\ &= \max \left(- \max_{m \neq j} \beta_{im}, q \right) \end{aligned} \quad (30)$$

1.2 Derivation of ψ_{ij}

$\psi_{ij}(h_{ii})$ is computed as:

$$\psi_{ij}(h_{ii}) = \max_{k \in \partial_j} \left(G_j(\partial_j) + \sum_{l \in \partial_j \setminus \{i\}} \phi_{lj}(h_{ll}) \right) \quad (31)$$

When $h_{ii} = 1$ $\Rightarrow h_{kk} = 0 \ \forall k \in \partial_j$. So:

$$\psi_{ij}(1) = \sum_{l \in \partial_j \setminus \{i\}} \phi_{lj}(0) \quad (32)$$

When $h_{ii} = 0$: There are two possibilities: either $h_{kk} = 0 \ \forall k \in \partial_j$ or one of $k \in \partial_j$ is 1 and rest are 0. We need to look for the possibility that maximizes the message $\psi_{ij}(0)$. So:

$$\begin{aligned} \psi_{ij}(0) &= \max \left(\sum_{l \in \partial_j \setminus \{i\}} \phi_{lj}(0), \max_{l \in \partial_j \setminus \{i\}} (\phi_{lj}(1) + \sum_{m \in \partial_j \setminus \{i, l\}} \phi_{mj}(0)) \right) \\ &= \max \left(0, \max_{l \in \partial_j \setminus \{i\}} (\phi_{lj}(1) - \phi_{lj}(0)) \right) + \sum_{l \in \partial_j \setminus \{i\}} \phi_{lj}(0) \\ &= \max \left(0, \max_{l \in \partial_j \setminus \{i\}} \phi_{lj} \right) + \sum_{l \in \partial_j \setminus \{i\}} \phi_{lj}(0) \end{aligned} \quad (33)$$

From (32) and (33), we get:

$$\psi_{ij} = \psi_{ij}(1) - \psi_{ij}(0) = - \max(0, \max_{l \in \partial_j \setminus \{i\}} \phi_{lj}) \quad (34)$$

1.3 Complexity Analysis of Messages

We now look at the equations from (18) to (23).

- β_{ij} : From Eq. (18), it is evident that non-diagonal messages are simply a summation of s_{ij} and α_{ij} . For diagonal messages, we also need sum over the involved neighborhoods \mathcal{M}_i . In worst case \mathcal{M}_i can have $n - 1$ points indicating that for diagonal, β_{ii} can be of $O(n)$. Hence computation of all β_{ii} messages is of $O(n^2)$ in worst case. Since non-diagonal

1277 messages computation is also of $O(n^2)$, we can state that
 1278 computation of all β_{ij} is $O(n^2)$.

- 1279 • ϕ_{ij} : Using the same argument as in case of β_{ij} we can state
 1280 that these messages are also of $O(n^2)$.
- 1281 • ρ_{ij} : These messages are also $O(n^2)$. The argument is the
 1282 same as for β_{ij} and ϕ_{ij} .
- 1283 • $\bar{\eta}_{ij}$: The computation of original messages η_{ij} of AP is of
 1284 $O(n^2)$ [13]. Since $\bar{\eta}_{ij}$ is just a maximum with a scalar, we can
 1285 state that these messages are also $O(n^2)$.
- 1286 • ϕ_{ij} : These messages are computed only for diagonal entries
 1287 h_{ij} . Each message can be $O(n)$ in worst case. So computation
 1288 of all messages is $O(n^2)$.
- 1289 • α_{ij} : These messages are the same as in AP since we did not
 1290 change the consistency constraint of AP. So computation of
 1291 these messages is also $O(n^2)$.

1292 Since all involved messages can be computed in $O(n^2)$ time, we
 1293 can state that the worst time complexity of RAP is $O(n^2)$. In practice
 1294 it reduces to $O(|\mathcal{E}|)$ where \mathcal{E} is the set of edges along which $s(i, j)$
 1295 values are defined.

2 ADDITIONAL FEATURES OF RAP

Besides the features that we already mentioned in Sec. 5, we will elaborate on some additional features of RAP in this section.

2.1 Pruning for Refined Inter-exemplar Connections:

Sec. 5.3 briefs how the inter-exemplar connections histogram can be used to analyse cluster strength in different regions of a cluster. An analyst can manually look at the weakly connected pairs to decide whether the links should be retained or *pruned*. We can go one step further to automate this process of pruning. Once RAP messages have converged, we select m nearest exemplars for every exemplar because we are interested in only the connections between nearby exemplars. For these $|\mathcal{E}|m$ exemplar pairs, we can simply prune all the boundary connections where connection strength is less than a predefined threshold n_t . The computational complexity of this process is $O(nm|\mathcal{E}|)$ which is less than $O(n^2)$ for all practical purposes. This step is performed during the decision phase before finding the connected components in Alg. 1. Hence, if needed, it can even be iterated for multiple values of n_t without having to run the message passing phase every time. In the example of Fig. 6 in our work, pruning with $n_t = 2$ separates all exemplar pairs with single boundary connection, thus breaking the single link in the weak strength region of the red cluster and resulting in 3 clusters. However, for 19 out of 20 exemplar pairs, both local exemplars belonging to the pair still fall in the same cluster. Only one exemplar pair now has its members in different clusters. Hence, by comparing the results of RAP with $n_t = 1$ (i.e. no pruning) and $n_t = 2$, this one pair can be identified programmatically without going through all 20 pairs manually. Therefore, by running the pruning process for different values of n_t , we not only split the discovered clusters at the weaker regions, but also find the data points that were responsible for merging these regions. In our work, we fix $n_t = 3$ independent of the dataset and other settings.

2.2 Local Exemplars Connected to a Point

When a data point is connected to more than one well separated local exemplars, this signifies that the data point shares a mix of properties of these local exemplars. The local exemplars connected to x_i are given in $\mathcal{L}[i]$ which can be computed in $O(|\mathcal{E}|)$. Hence we can find the number of local exemplars connected to each data point in $O(n|\mathcal{E}|)$. This information can be visualized using a histogram where the x-axis represents the number of local exemplars and the height of each bar represents the number of data points connected to these many local exemplars. Such a histogram can be used for various purposes. It can allow for efficient hyperparameter tuning which we will discuss in Sec. 3. It can also be used to recognize the relative densities of different clusters in a dataset. We illustrate this on different datasets shown in Fig. 10. For this purpose we choose $\epsilon > 100$ percentile of $s(i, j)$, i.e., $G_j(\cdot)$ constraints are rendered inactive. As mentioned in Sec. 4.4 this often has a negligible impact on the global cluster discovery. At the same time this setting encourages multiple local exemplars to appear close by in dense regions, allowing a better analysis of relative densities.

Fig. 10(a) shows that for *Blobs* toy dataset [21] with clusters having equal densities, the histogram bars are co-located. On the other hand, the bars for red cluster in Fig. 10(b) are located further to the right of the bars for the green cluster, indicating that for red cluster the data points are connected to more local exemplars, which in turn implies a higher density of the red cluster. Similar behavior can be seen for the blue cluster in Fig. 10(c) and green cluster in Fig. 10(d).

2.3 Clustering New Data Points

The local exemplars, discovered by RAP, tend to lie in the dense regions of the clusters rather than the boundaries that may have noisy data. Therefore, compared to nearest neighbors, they serve as better reference points for analyzing cluster assignment confidence as well for clustering new points. Suppose that we want to cluster a new data point $x_{new} \notin \mathcal{X}$ based upon the clustering results of \mathcal{X} . For this purpose, we can compare x_{new} to the already found local exemplars (\mathcal{E}) and assign it to the cluster which contains the closest local exemplar. This usually lowers the complexity by orders of magnitude when compared to assigning a cluster by finding the closest neighbor(s) in the already clustered dataset, as would be the case for algorithms that do not provide any local information. The new data point x_{new} can be treated based upon its similarity with its potential exemplars.

- If x_{new} has a high enough similarity with two or more local exemplars belonging to the same cluster, connecting x_{new} to all of them will not affect cluster assignments.
- If x_{new} has a high enough similarity with two or more local exemplars belonging to different clusters, the case may be forwarded to a human analyst for expert opinion in order to decide whether x_{new} should be assigned to one of the two clusters or should the clusters be merged.
- If the new data point has low enough similarity to all of the local exemplars, it can be treated as an outlier.

Fig. 11 illustrates an example of clustering new data points where already discovered local exemplars are circled in blue color and new points to be clustered are bordered in red and green colors. We

1335
 1336
 1337
 1338
 1339
 1340
 1341
 1342
 1343
 1344
 1345
 1346
 1347
 1348
 1349
 1350
 1351
 1352
 1353
 1354
 1355
 1356
 1357
 1358
 1359
 1360
 1361
 1362
 1363
 1364
 1365
 1366
 1367
 1368
 1369
 1370
 1371
 1372
 1373
 1374
 1375
 1376
 1377
 1378
 1379
 1380
 1381
 1382
 1383
 1384
 1385
 1386
 1387
 1388
 1389
 1390
 1391
 1392




Figure 10: Relative Densities with right column showing RAP results for $\epsilon > 100$ percentile of $s(i, j)$. Left column shows the relative cluster density histograms. For (c) and (d) same parameters are used as mentioned in Sec. 6.1. The figure shows displaced histogram bars for clusters with different relative densities. Best viewed in high resolution.

can observe that the red point is most similar to the bottom-right exemplar while the green point is visually similar to both top-left and center-left exemplars. Therefore based upon the given pairwise similarities, we may choose one or multiple local exemplars to represent the new data point. We can apply this procedure to multiple new data points as well. In that case we should just be more careful that if we notice a sufficiently large number of new data points being declared as outliers this may be an indication of a new cluster being formed which was not present in the clustering results for X using RAP. Also if we notice enough data points having close enough similarities to two exemplars from different




Figure 11: RAP results sample on MNIST with local exemplars encircled in blue. We cluster new data points, bordered in red and green, by assigning them to the most similar exemplars.

clusters, they may form a strong enough bridge between the two clusters to merge them to one. In these cases it is a good idea to either re-cluster the whole dataset or a partial subset of the dataset.

2.4 Outliers and Noisy Data

We briefly mentioned for in Sec. 6.1 how RAP handles noisy boundaries between clusters for different synthetic datasets. We will now explore how RAP handles outliers. It starts by considering all data points as potential local exemplars and then explores the neighborhood of each data point to later decide if a data point should become a local exemplar or not. This provides RAP the ability to discover outliers as they do not tend to connect to other local exemplars that belong to well-formed bigger far away clusters. This is depicted in Fig. 12, where we have the half-moons dataset with 2 clusters and we have now added some outliers to it. Specifically, we have a lone outlier on the top right side of the figure and a small cluster of three outliers on the bottom left of the figure. RAP recognizes the lone outlier as a single point cluster, represented by the purple color. Similarly, RAP also puts together the three outliers on the lower-left corner as one small cluster, designated by blue color. It is also important to note that the presence of these outliers does not impact the clustering of the two well-formed bigger clusters. This can be observed by comparing Fig. 12 to Fig. 3(b), where we have the same dataset but without the outliers.

3 HYPERPARAMETERS - TUNING AND SENSITIVITY

3.1 Successive Tuning Methodology

By keeping ϵ fixed to a small value e.g. 99 percentile of input $s(i, j)$ values, we are left with two parameters, p and q . We can perform efficient successive tuning of the hyperparameters easily to obtain desired results in a few steps. This is due to the following reasons:




Figure 12: Half-moons dataset with outliers. These results are obtained using the same parameters as Fig. 3(b), i.e., $p = 0.6$, $q = -0.97$ and $\epsilon = 0.99$.

- It is easy and intuitive to understand the impact of changing any of the two hyperparameters on the clustering outcome.
- The abundance of insightful information, beyond just cluster assignments, provided by RAP at each step helps the analyst in deciding how to adapt the hyperparameters to improve the results.

Some ways to utilize local information to guide the subsequent steps in hyperparameter tuning are as follows:

- **Too few local exemplars:** A small number of local exemplars indicates that we have set p too low and we should increase it such that we obtain enough local exemplars that are spread throughout the dataset.
- **Boundary connections:** Once we have an adequate number of local exemplars available, we focus on building boundary connections between the local exemplars for global structure discovery. For this we use inter-exemplar connections histogram to tune q as follows
 - **Inter-exemplar connections histogram shifted too much to the left:** This configuration implies that the maximum linkage penalty is too high and hence not allowing sufficient boundary connections to form between the local exemplars. This will also manifest itself in terms of many more clusters than expected. In this scenario, we should decrease the linkage penalty so that more boundary connections can appear.
 - **Inter-exemplar connections histogram shifted too much to the right:** This scenario is the opposite of the previous one. Due to the low linkage penalty data points also get connected to far off local exemplars (which possibly may belong to a different “ground truth” cluster). This can result in too few clusters than expected and violates the basic principle of introducing the penalty q that we want to form only boundary connections between close enough local exemplars to facilitate global cluster

discovery. Allowing points to connect to even far off local exemplars rather distorts the global cluster discovery by even merging well-separated clusters. In this scenario, we should increase the linkage penalty.

- **Confidence values:** The confidence values discussed in Sec. 5.2 can also be used to determine if the results obtained for a specific choice of hyperparameters are satisfactory. Analysis of a few points with low confidence values can point out whether we have two clusters lying close by, or that the algorithm has not been able to form boundary connections to merge sub-clusters where it should have. In the second case we need to reduce the linkage penalty to merge these sub-clusters.

Finally, we want to emphasize that all the results presented in this work, unless explicitly stated otherwise, are obtained via successive tuning of hyperparameters rather than any kind of parameter sweeps. Hence these results are a good representation of what one can expect in the normal application of RAP and do not correspond to the ceiling performance on the respective datasets. Furthermore, we also notice that the final results shown in Sec. 6.1 and Sec. 6.2 correspond to hyperparameter values that are very similar for different datasets although the datasets have very different characteristics. This also points to the ease of finding suitable hyperparameters for a wide variety of datasets with different cluster characteristics by exploring a very narrow range of hyperparameter values, as long as the datasets employ the same (or similar) pairwise similarity metric.

3.1.1 Successive Tuning Example. We will now present a step by step example of using the insights from Sec. 4 and Sec. 3.1 to do hyperparameter tuning for the half-moons dataset earlier used in Fig. 1 and Fig. 7. We keep $\epsilon = 99$ percentile of $s(i, j)$. These steps are illustrated in Fig. 13 where each row corresponds to a step. The columns list clustering results with exemplars, inter-exemplar connection strength, exemplars connected to data points and cluster assignment confidence values respectively. Note that RAP only needs $s(i, j)$ values and the plots in the first and fourth columns are for sake of illustration only. The relevant information in these plots is also available otherwise for the analyst.

- (1) As the first step, we apply RAP to the dataset for $p = 60$ percentile of $s(i, j)$ while keeping $q = -\infty$. This gives us 60 clusters as shown in the first row of Fig. 13. The results show that there is only one exemplar connection per point as we have set the penalty q to maximum. As a result there is no need for inter-exemplar links as there is no exemplar pair with common points. The analysis of cluster assignment confidence shows that there are many similar points that have been assigned to different clusters. These points can easily be obtained programmatically for an analyst to decide whether they should be connected or not. We accept this relatively large number of clusters in the first step because we know that they will be merged in subsequent steps. If required, we can also check the dependence of results on p by playing around with this parameter. For instance, the current dataset yields 73 and 38 clusters for $p = 80$ and $p = 20$ percentile of $s(i, j)$ respectively.




Figure 13: Successive tuning of hyperparameters p and q for half-moons dataset. We use the insights provided by current step to refine our choice of p and q for next step.

- (2) First step shows that there is a need to connect many closely located clusters. For second step, we fix $p = 60$ percentile of $s(i, j)$ and select $q = 95$ percentile of $s(i, j)$ to encourage multiple exemplars per data point. However this results in a single cluster. We can see from the histograms of second row in Fig. 13 that on average 5 points are connected to a point. Moreover, inter-exemplar connections are quite high. There are exemplar pairs with more than 60 common points. This indicates that we need to increase the penalty to force lesser common connections between exemplars. There is no need for a confidence plot here as we have only one cluster.
- (3) Based upon insights from the first and the second step, we now need to have q somewhere between 100 and 95 percentile of $s(i, j)$ while keeping p fixed. We choose $q = 97.5$ (i.e. mean of 100 and 95) percentile and get 2 clusters as shown in the third row of Fig. 13. The histograms now show that most data points are now connected to 1, 2 or 3 local exemplars which is a reasonable result. Moreover, the inter-exemplar connections are also reduced compared to the second step, although there are still some pairs with more than 22 common points. It now depends upon the analyst if he wants to further reduce these connections or if he is content with the results. The co-located bars in the histogram of exemplar connections per point also indicate that the two discovered clusters have similar density.

3.2 Sensitivity of Clusters Discovery to Hyperparameters

We look at how global structure discovery is impacted by the variation of individual hyperparameters. For this purpose, we use the example in Sec. 3.1.1 and start with the hyperparameters set selected in the last step of successive tuning i.e. $p = 60$, $q = -97.5$ and $\epsilon = 99$ percentile of $s(i, j)$ values. We then change one parameter at a time to see the sensitivity of the results.

3.2.1 Sensitivity to p : For fixed value of q and ϵ , p controls the number of local exemplars. In general as long as p is not too low the global cluster results don't get impacted as shown in Fig. 14. This is unlike AP where variations in p have a strong impact on the number of clusters obtained and the cluster assignments. In the case of RAP, as long as p is varied in a reasonable range, this only changes the number of local exemplars appearing in any region but as long as the number is not too small and they can get connected via boundary connections, the number of globally discovered clusters does not change. Only when p is too low, we are no longer able to form bridge connections between the appropriate exemplars and we lose the ability to discover the global structure.

3.2.2 Sensitivity to q : q is arguably the hyperparameter which needs the most care but there is a reasonable range where the global clusters obtained are not affected and successive tuning can discover a suitable value of q in this range quickly. When we have




Figure 14: Sensitivity to p while keeping q and ϵ constant. All values are in percentiles of input $s(i, j)$ values. Best viewed in high resolution




Figure 15: Sensitivity to q while keeping p and ϵ constant. All values are in percentiles of input $s(i, j)$ values. Best viewed in high resolution




Figure 16: Sensitivity to ϵ while keeping p and q constant. All values are in percentiles of input $s(i, j)$ values. Best viewed in high resolution

a too high linkage penalty, clusters start to break at their weakly connected regions as seen in Fig. 15. Note that still in Fig. 15 there have not been drastic changes in the results beyond the green cluster being broken where it was most weakly connected.

3.2.3 Sensitivity to ϵ : Fig. 16 shows the effect of changing ϵ on global clusters discovery while keeping p and q fixed. It can be seen that for small values of ϵ , the global clusters discovery is not affected. However, if the purpose of ϵ , which is only to avoid crowding of local exemplars in dense regions, is not kept in mind and the values are changed beyond a limit, they start affecting the clusters. This can be seen in the case when $\epsilon = 98$. It is worth noticing that again, the cluster gets split at the weakly connected region. Practically, ϵ can be considered as a fixed parameter at 99 percentile of input $s(i, j)$.

Dataset	Points	True clusters
Aggregation	788	7
Flame	240	2
R15	600	15
Circles	1500	2
Optdigits	1797	10
MNIST	1000	10
Skin	5000	2
PPI	1171	276

Table 4: Statistics of the synthetic and real-world datasets

4 EXPERIMENT DETAILS

4.1 Datasets

Tab. 4 gives the number of data points and number of ground truth clusters for all the datasets used in the experiments. For synthetic

1857 datasets, these are the only relevant statistics as all the synthetic
 1858 datasets are in \mathbb{R}^2 and the plots have already been given in Fig. 7.
 1859 The real-world datasets are further explained below

- 1860 • **Optdigits** consists of grey scale 8×8 images of handwritten
 1861 digits from 0 to 9. We only use the test set, consisting of 1797
 1862 images, for clustering.
- 1863 • **MNIST** is also a similar dataset of 28×28 gray scale images of
 1864 handwritten digits. We choose 1000 images from the dataset
 1865 at random to define X .
- 1866 • **Skin** dataset has of 245,057 points constructed over (B, G,
 1867 R) color space, consisting of two classes: skin tones and non-
 1868 skin tones. The skin tones are obtained using skin textures
 1869 from face images of people of different ages, gender and
 1870 ethnicity. We randomly choose a subset of 5000 points that
 1871 are equally distributed between the two classes.
- 1872 • **Protein-Protein Interactions** is the Saccharomyces Cere-
 1873 visiae gold standard dataset. The pairwise similarity values
 1874 for this dataset represent protein-protein interactions. We
 1875 use the probabilistic measure, proposed in [24], to define our
 1876 pairwise similarity matrix. This dataset is different from the
 1877 previously discussed synthetic and real datasets in the sense
 1878 that the available ground truth does not necessarily assign
 1879 each data point to only one ground truth cluster. Further-
 1880 more, some ground truth clusters are a subset of other ground
 1881 truth clusters. To reduce such overlap in the discovered clus-
 1882 ters, we threshold the protein-protein interactions such that
 1883 the interactions below 0.4 are set to 0². Furthermore, we
 1884 only consider proteins that are common with CYC2008. This
 1885 leaves us with 1171 proteins (data points) assigned to 276
 1886 ground truth clusters.

1888 4.2 Implementation Settings

1889 For AP, we report the performance using the default parameter
 1890 selection ($p = 50$ percentile of input $s(i, j)$) as well as the ceiling
 1891 performance obtained by sweeping p from 10 to 90 percentile of
 1892 input similarities with a percentile step size of 5. All the algorithms,
 1893 except AP, require tuning of two hyperparameters. For HAP and
 1894 SCAP, although there exist good intuitions for the hyperparameters
 1895 involved, there is no simple way of choosing them. Therefore we
 1896 tune the parameters of HAP and SCAP manually by searching over
 1897 preference ranges between 10 and 80 percentile of $s(i, j)$ values
 1898 with a percentile step size of 5 and present the average accuracy
 1899 score of the top 3 results. For MEAP, we use the default settings of
 1900 parameters as suggested by the authors [30] i.e. the preference and
 1901 linkage parameters are set to the medians of pairwise similarities
 1902 and linkage matrix respectively, where linkage matrix is computed
 1903 by normalizing pairwise similarity matrix by n . It is a common
 1904 practice in message-passing based algorithms to aid convergence
 1905 by *damping* every message m in iteration $(t + 1)$ by a factor of
 1906 $\lambda \in [0, 1]$ i.e.

$$1907 m^{t+1} := (1 - \lambda)m^{t+1} + \lambda m^t. \quad (35)$$

1908 The damping factor λ is kept at 0.8 in all experiments for all the
 1909 message passing based algorithms. It is worth noticing that within
 1910 reasonable ranges (e.g. $\lambda \geq 0.5$), the results remain the same and

1912 ²This approach has been adopted in previous studies too e.g. [29] and [22], although
 1913 different thresholds have been used there.

1915 λ only affects the speed of convergence. Moreover, maximum it-
 1916 erations are fixed to 1000 and convergence-iterations are fixed to
 1917 100. i.e. if the results remain the same for 100 iterations, the al-
 1918 gorithm stops irrespective of the number of maximum iterations.
 1919 We keep the number of maximum-iterations high to encourage all
 1920 the competitors to converge as some algorithms might take a bit
 1921 longer to converge compared to the others. For MCL, we sweep
 1922 inflation from 1 to 5 with step-size of 1., and report the best results.
 1923 For DBSCAN and OPTICS, we use the default parameter settings of
 1924 sklearn python package. The boundary points for DBSCAN and
 1925 OPTICS are clustered using 5 nearest neighbors.

1926 For RAP, in all the experiments, we used successive tuning for
 1927 hyperparameter selection, unless specified otherwise. We can see
 1928 that successive tuning may lead to a self-preference value higher
 1929 than the default for AP i.e. 50-percentile of $s(i, j)$. This is intuitive
 1930 because in RAP, we have a relatively higher number of the local
 1931 exemplars with sub-clusters that are subsequently merged using
 1932 boundary connections. All the messages are initialized to 0. For
 1933 inter-exemplar connections histogram, we keep $m = 5$ i.e. we
 1934 target only 5 nearest exemplars for every exemplar. In addition to
 1935 the heuristic performance, we also present the ceiling performance
 1936 obtained by sweeping p over the percentile range from 10 to 80
 1937 percentile of $s(i, j)$ with the percentile step size of 5, and q from
 1938 90 to 99 percentile of $s(i, j)$ with the percentile step size of 1. Rest
 1939 of the settings are kept the same as other message passing based
 1940 competitors. Implementation is available at [3].